

Supplementary Appendix for *Delegation Space-Saving: Fast, Accurate, and Concurrent Frequent Elements Queries and Updates*

Anonymous Author(s)

ACM Reference Format:

Anonymous Author(s). 2020. Supplementary Appendix for *Delegation Space-Saving: Fast, Accurate, and Concurrent Frequent Elements Queries and Updates*. In *51st International Conference on Parallel Processing (ICPP 2022)* August 29 - Sept 1, 2022 Bordeaux, France. ACM, New York, NY, USA, 3 pages. <https://doi.org/tobefined>

1 APPENDIX

Due to space limitations, below we include some complementary material that could not fit in the 10-page limit:

Proof of lemma 1

PROOF. Let ss be a Space-Saving algorithm instance that maintains $k = \frac{1}{\epsilon}$ counters and that processes the stream S , and let ss_i be one of T Space-Saving algorithm instances that each processes the sub-stream S_i . The established guarantee of ss [2] is to report \mathcal{F} by tracking all elements $e \in U$ such that

$$f(e) > N\epsilon = \frac{N}{k} \quad (1)$$

We have that the length of stream S_i is equal to $N_i = \frac{N}{T}$ due to domain splitting. An algorithm instance ss_i that maintains k counters will then report \mathcal{F}_i , the ϵ -approximate frequent elements that are part of U_i , by tracking all elements $e \in U_i$ such that

$$f(e) > \frac{N}{Tk} \quad (2)$$

If we reduce the number of counters maintained by ss_i to $\frac{k}{T}$, then ss_i is guaranteed to track all elements e such that

$$f(e) > \frac{N}{T \frac{k}{T}} = \frac{N}{k} = N\epsilon \quad (3)$$

Since (1) and (3) are the same, we can conclude that each of the T different ss_i instances can keep just $\frac{k}{T} = \frac{1}{\epsilon T}$ counters to have the same accuracy guarantee as ss , given that the union of each \mathcal{F}_i is output by Delegation Space-Saving. \square

Proof of lemma 2

PROOF. Here we slightly modify the proof of Theorem 5 in [2] (inequality 8 and onward).

Let $k = \frac{1}{\epsilon}$ denote the number of counters and let $H_{n,a} = \sum_{i=k+1}^n \frac{1}{i^a}$.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ICPP 2022, August 29 – Sept 1, 2022, Bordeaux, France

© 2020 Association for Computing Machinery.

ACM ISBN to be defined...\$15.00

<https://doi.org/tobefined>

Lemma 3.4 and Equation (1) from [2] state that the maximum least value of a counter in Space-Saving is $\min_{max} \geq \frac{N-F_k}{k}$, where F_k is the sum of occurrences of the k most frequently occurring elements in S . Any element e with $f(e) > \min_{max}$ is guaranteed to be monitored by Space-Saving, regardless of stream permutation (lemma 4.3 [2]). We have that $N - F_k$ is equal to the sum of occurrences of elements ranked $k+1$ through $|U|$. An element e has occurred $f(e) = \frac{N}{H_{|U|,a} r(e)^a}$ times in a stream of length N constructed by drawing elements from the Zipf distribution with parameter a . We can encode this reasoning in inequality 4

$$f(e) > \min_{max} \stackrel{zipf}{=} \frac{1}{r(e)^a} \frac{N}{H_{|U|,a}} > \frac{1}{k} \frac{N}{H_{|U|,a}} \sum_{i=k+1}^{|U|} \frac{1}{i^a} \quad (4)$$

Simplifying yields:

$$\frac{1}{r(e)^a} > \frac{1}{k} \sum_{i=k+1}^{|U|} \frac{1}{i^a} \quad (5)$$

Metwally et al. then approximates the expression in the right-hand side of inequality (5):

$$\frac{1}{r(e)^a} > \frac{1}{k^a} \sum_{i=2}^{\frac{|U|}{k}} \frac{1}{i^a} \quad (6)$$

$\sum_{i=2}^{\frac{|U|}{k}} \frac{1}{i^a} < \zeta(a) - 1$, so we can let k satisfy a stronger constraint:

$$\frac{1}{r(e)^a} > \frac{1}{k^a} (\zeta(a) - 1) \quad (7)$$

We continue by considering an assumption that the sum of all occurrences of the non-frequent elements represented by $(\zeta(a) - 1)$ are evenly distributed to all threads due to domain splitting. (This assumption is motivated by the fact that $N - F_k$ tends to be much smaller than F_k , and at the same time, $|U| - (k + 1) \gg k$, in the case that the input follows a Zipf distribution with $a > 1$.)

$$\frac{1}{r(e)^a} > \frac{1}{k^a} \frac{(\zeta(a) - 1)}{T} \quad (8)$$

This simplifies to:

$$k > r(e) \left(\frac{\zeta(a) - 1}{T} \right)^{\frac{1}{a}} \quad (9)$$

We now direct our attention to the inequality $N\epsilon > \min_{max} > \frac{N}{r(e)^a \zeta(a)}$, which describes the least ranked element that exceeds ϵN . It can be solved for $r(e)$:

$$r(e) < \left(\frac{1}{\epsilon \zeta(a)} \right)^{\frac{1}{a}} \quad (10)$$

We can now substitute r in equation 9:

$$k > \left(\frac{\zeta(a) - 1}{T\epsilon\zeta(a)} \right)^{\frac{1}{a}} = \left(\frac{1}{T\epsilon} - \frac{1}{T\epsilon\zeta(a)} \right)^{\frac{1}{a}} \quad (11)$$

Thus, setting $k = \left(\frac{1}{T\epsilon} \right)^{\frac{1}{a}}$ guarantees that inequality (4) is satisfied, i.e. all elements with $f(e) > \epsilon N$ are *monitored* (and hence might be delivered) by the Space-Saving instance. Therefore, using $k = \left(\frac{1}{T\epsilon} \right)^{\frac{1}{a}}$ counters at each Space-Saving instance in Delegation Space-Saving algorithm solves the ϵ -approximate frequent elements problem, given that all algorithm instances are considered when reporting the frequent elements. \square

Proof of lemma 5

PROOF. If we consider an unbounded universe of possible elements to target the general case, then the probability of drawing element e is $\frac{1}{\zeta(a)r(e)^a}$, where $r(e)$ denotes the rank of element e . Since there are at most mT counts of an element in Delegation Filters kept for the owner of e by lemma 4, then since these mT elements were drawn from a noiseless Zipf distribution, there can only be at most $\frac{mT}{\zeta(a)r(e)^a}$ counts of element e . \square

Proof of claim 1

PROOF. We have that $f_{N_S}(e) \leq f_{N_E}(e)$. There are at most $\mathbb{D}(e)$ counts of element e that have not yet been inserted into the Space-Saving instance of the owner of e , therefore $f_{N_S}(e) - \mathbb{D}(e) \leq \hat{f}_N(e)$ is the minimum value a counter can assume. The maximum over-estimation of Space-Saving is ϵN , which is maximized at the end of a query, when N_E elements have been processed. Therefore the estimated count of an element is at most $\hat{f}_N(e) \leq f_{N_E}(e) + \epsilon N_E$. \square

Proof of claim 2

PROOF. The Delegation Space-Saving algorithm reports all elements with estimated count

$$f(e)_{N_S} > \phi N_S \quad (12)$$

This is true since N_S is calculated at the start of a query, all elements with $\hat{f}_N(e) > \phi N_S$ are reported, and Space-Saving counters increase monotonically.

As shown in claim 2, $\hat{f}(e)$ is at least (a) $f(e)_{N_S} - \mathbb{D}(e)$ and at most (b) $f(e)_{N_E} + N_E\epsilon$.

Substituting $\hat{f}(e)$ for (a) in (12) yields:

$$f(e)_{N_S} > \phi N_S + \mathbb{D}(e) \quad (13)$$

Substituting $\hat{f}(e)$ for (b) in (12) yields:

$$f(e)_{N_E} > \phi N_S - \epsilon N_E \quad (14)$$

Expression (a) and the negation of (b) together give that all elements $f(e)_{N_S} > \phi N_S$ and no elements $f(e)_{N_E} < \phi N_S - \epsilon N_E$ are reported by the Delegation Space-Saving algorithm. \square

Proof of claim 3

PROOF. We normalize the expression from inequality (13):

$$\frac{f(e)}{N_S} > \phi + \frac{\mathbb{D}(e)}{N_S} \quad (15)$$

$\mathbb{P}(e) = \frac{f_{N_S}(e)}{N_S}$ is the probability at which element e occurs in the underlying input distribution. We then have that:

$$\lim_{N_S \rightarrow \infty} \mathbb{P}(e) > \phi + \frac{\mathbb{D}(e)}{N_S} \rightarrow \mathbb{P}(e) > \phi \quad (16)$$

Meaning that all elements with $\mathbb{P}(e) > \phi$ are reported given that enough elements have been processed. \square

Proof of claim 4

PROOF. Using the expression from inequality (13), we can rewrite it in its Zipfian form:

$$\frac{1}{\zeta(a)r(e)^a} > \phi + \frac{mT}{\zeta(a)r(e)^a N_S} \quad (17)$$

Solving for $r(e)$ gives:

$$\left(\frac{1}{\zeta(a)\phi} - \frac{mT}{\zeta(a)\phi N_S} \right)^{\frac{1}{a}} > r(e) \quad (18)$$

Simplifying yields:


$$\left(\frac{1}{\zeta(a)\phi} \right)^{\frac{1}{a}} \left(1 - \frac{mT}{N_S} \right)^{\frac{1}{a}} > r(e) \quad (19)$$

We then have that:


$$\lim_{N_S \rightarrow \infty} \left(\frac{1}{\zeta(a)\phi} \right)^{\frac{1}{a}} \left(1 - \frac{mT}{N_S} \right)^{\frac{1}{a}} > r(e) \rightarrow \left(\frac{1}{\zeta(a)\phi} \right)^{\frac{1}{a}} > r(e) \quad (20)$$

Meaning that all elements with a rank lower than $\left(\frac{1}{\zeta(a)\phi} \right)^{\frac{1}{a}}$ are guaranteed to be reported given that enough elements have been processed. \square


Scalability results, CAIDA direction B



(a) 0% queries.



(b) 0.01% queries.



(c) 0.02% queries.

Figure 1: Dir B scalability

REFERENCES

- [1] Anonymous. Delegation Space-Saving. <https://github.com/anonymous7495/Delegation-Space-Saving/tree/main>, 4 2022.
- [2] A. Metwally, D. Agrawal, and A. E. Abbadi. An integrated efficient solution for computing frequent and top- k elements in data streams. *ACM Transactions on Database Systems*, 31(3):1095–1133, 2006.

Reproducibility: The code is available for reproducibility purposes: [1].